



3.5.0

Overview

Programming Guides ▾

API Docs ▾

Deploying ▾

More ▾

Search the docs

Structured Streaming integration for Kafka 0.10 to read data from and write data to Kafka.

Linking

For Scala/Java applications using SBT/Maven project definitions, link your application with the following artifact:

```
groupId = org.apache.spark  
artifactId = spark-sql-kafka-0-10_2.12  
version = 3.5.0
```

Please note that to use the headers functionality, your Kafka client version should be version 0.11.0.0 or up.

For Python applications, you need to add this above library and its dependencies when deploying your application. See the [Deploying](#) subsection below.

For experimenting on `spark-shell`, you need to add this above library and its dependencies too when invoking `spark-shell`. Also, see the [Deploying](#) subsection below.

Reading Data from Kafka

Creating a Kafka Source for Streaming Queries



Data Science & Engi...

Create

Workspace

Recents

Search

Data

Compute

Workflows

Menu options

Compute >

my cluster

Configuration

Filter lib...

Status

Install library [Provide feedback](#)

Library Source

☐ Upload☐ DBFS/S3☐ PyPI☒ Maven☐ CRAN☐ Workspace Library

Coordinates

org.apache.spark:spark-sql-kafka-0-10_2.12:3.5.0

[Search Packages](#)

Repository ?

Optional

Exclusions

Dependencies to exclude (log4j:log4j;unit:junit)

Cancel

Install

More ...

Terminate

Edit

compute UI - Master

Uninstall

Install new



Cmd 1

```
BOOTSTRAP_SERVER = "pkc-6ojv2.us-west4.gcp.confluent.cloud:9092"  
JAAS_MODULE = "org.apache.kafka.common.security.plain.PlainLoginModule"  
CLUSTER_API_KEY = "TQ7F3730G6J4TT4G"  
CLUSTER_API_SECRET = "ftb0mWOYcCT6oAG9GSIjyUxHYuJmbUPDi0G0c5ymt0t6FKJ1J1IVDbDRVqKoG5H4"
```

Cmd 2

Python ▶ ▾ - x

Shift+Enter to run

Shift+Ctrl+Enter to run selected text



09-kafka-consumer Python

File Edit View Run Help Last edit was 7 minutes ago Provide feedback

Run all

Connect

Share

Publish

Cmd 1

```
BOOTSTRAP_SERVER = "pkc-6ojv2.us-west4.gcp.confluent.cloud:9092"
JAAS_MODULE = "org.apache.kafka.common.security.plain.PlainLoginModule"
CLUSTER_API_KEY = "TQ7F3730G6J4TT4G"
CLUSTER_API_SECRET = "ftb0mWOYcCT6oAG9GSIjyUxHYuJmbUPD10G0c5ymt0t6FKJ1J1IVDbDRVqKoG5H4"
```

Cmd 2

```
df = ( spark.read
      .format("kafka")
      .option("kafka.bootstrap.servers", BOOTSTRAP_SERVER)
      .option("kafka.security.protocol", "SASL_SSL")
      .option("kafka.sasl.mechanism", "PLAIN")
      .option("kafka.sasl.jaas.config", f"{JAAS_MODULE} required username='{CLUSTER_API_KEY}' password='{CLUSTER_API_SECRET}';")
      .option("subscribe", "invoices")
      .load()
    )
```

Cmd 3

```
display(df)
```



09-kafka-consumer Python

File Edit View Run Help Last edit was 9 minutes ago Provide feedback

Run all

my cluster

Share

Publish

(2) Spark Jobs

Table +

	key	value
1	U1RSODUxMw==	eyJbnZvaWNITnVtYmVyljogljcxMzcyOTczliwglkNyZWFOZWRUaW1lIjogMTU5NTY4ODkwMTIxOSwglIN0b3JISUQiOiMylsICJQb3NJRCI6ICJQT1MxNjMiLCAiQ2FzaGlcklEljoglk9BUzk2MSIsICJdDdXM= (truncated)
2	U1RSODUxMw==	eyJbnZvaWNITnVtYmVyljogljcxNjQ3NTEzliwglkNyZWFOZWRUaW1lIjogMTU5NTY4ODkwMTIxOSwglIN0b3JISUQiOiMylsICJQb3NJRCI6ICJQT1M4ODEiLCAiQ2FzaGlcklEljoglk9BUzM1NCIsICJdDdXM= (truncated)
3	U1RSMjYyOQ==	eyJbnZvaWNITnVtYmVyljogljQ4ODg2MDciLCAiQ3JlYXRlZFRpbWUiOiAxNTk1Njg4OTAzMjU4LCAiU3Rvcml6I6ICJdDdXM= (truncated)
4	U1RSNjE2Mg==	eyJbnZvaWNITnVtYmVyljogljQ4ODg2MDciLCAiQ3JlYXRlZFRpbWUiOiAxNTk1Njg4OTAzMjU4LCAiU3Rvcml6I6ICJdDdXM= (truncated)
5	U1RSODUxMw==	eyJbnZvaWNITnVtYmVyljogljQ4MzAyNjQ1liwglkNyZWFOZWRUaW1lIjogMTU5NTY4ODkwMTIxOSwglIN0b3JISUQiOiMylsICJQb3NJRCI6ICJQT1MxNTUiLCAiQ2FzaGlcklEljoglk9BUzU1OCIsICJdDdXM= (truncated)
6	U1RSODUxMw==	eyJbnZvaWNITnVtYmVyljogljcxMzcyOTczliwglkNyZWFOZWRUaW1lIjogMTU5NTY4ODkwMTIxOSwglIN0b3JISUQiOiMylsICJQb3NJRCI6ICJQT1MxNjMiLCAiQ2FzaGlcklEljoglk9BUzk2MSIsICJdDdXM= (truncated)

30 rows | 13.30 seconds runtime

Command took 13.30 seconds -- by prashantkumarpandey@gmail.com at 10/22/2023, 3:05:07 PM on my cluster

Shift+Enter to run

Shift+Ctrl+Enter to run selected text



edemv

▶ (2) Spark Jobs

Table ▾ +

		topic	partition	offset	timestamp	timestampType
1	SUQioAiU1RSODUx	invoices	1	0	2023-10-21T10:31:56.031+0000	0
2	SUQioAiU1RSODUx	invoices	1	1	2023-10-21T10:31:57.036+0000	0
3	RCI6ICJTVFlyNjI5liwgl	invoices	1	2	2023-10-21T10:31:59.043+0000	0
4	3JISUQioAiU1RSNjE	invoices	1	3	2023-10-21T10:31:59.360+0000	0
5	SUQioAiU1RSODUx	invoices	1	4	2023-10-21T10:31:59.610+0000	0
6	SUQioAiU1RSODUx	invoices	1	5	2023-10-21T10:56:26.027+0000	0

↓ 30 rows | 13.30 seconds runtime

Command took 13.30 seconds -- by prashantkumarpandey@gmail.com at 10/22/2023, 3:05:07 PM on my cluster

Shift+Enter to run

Shift+Ctrl+Enter to run selected text

